

Principal Components Analysis

Dimensionality

Let's start thinking about our data in terms of *dimensions in space*.

Each **predictor** is an **axis**.

The values of the predictors for a certain **observation** define a point in space.

When we compute distances between observations for KNN, we are computing **distance in space**.

When we fit a **regression**, we are drawing a **straight line** through the points.

The curse of dimensionality

We run into trouble when we have **too many dimensions**.

What does "too many" mean?

Parametric estimation -> We can't estimate 7000 coefficients from only 44 observations!

Interpretability -> Do we really want to translate our model into meaning for thousands of predictors???

Flexibility -> More predictors = more flexibility = overfitting?

Principal Components Analysis

PCA is a way to *transform our data* (prior to modeling!) so that it has fewer **dimensions in space**.

Instead of:

axis 1 = Predictor A

axis 2 = Predictor B


axis 3 = Predictor C

axis 1 = 0.5 (Pred A) + 0.2 (Pred B) + 0.3 (Pred C)

axis 2 = 0.1 (Pred A) + 0.7 (Pred B) + 0.2 (Pred C)

axis 3 = 0.1 (Pred A) + 0.2 (Pred B) + 0.8 (Pred C)

PCA



PCA

1. **Standardize** all axes.
2. Find the axis of **highest variance**: This is PC 1.
3. Find the axis of **highest variance** that is **perpendicular to PC 1**: This is PC2.
4. Continue until you have p PCs, where p = number of predictors (or, if $p > n$, until you have n PCs).
5. Use only the first k predictors in your analysis, where $k < p$ and $k < n$.


Example

The **Federalist papers** are a series of essays written by John Jay, Alexander Hamilton, and James Madison.

Data: How many of each word was used in each essay (for the most common 200 words only).

Example

If we choose a couple words and plot our data...



Example

Instead, let's apply pca:

```
fed_matrix <- fed_ex %>% select(-auths_known) %>% as.matrix()  
pc <- prcomp(fed_matrix, center = TRUE, scale = TRUE)
```

Example

Combinations of variables that create new axes:

##	PC1	PC2	PC3	PC4	PC5	PC6
## a	0.1981759	-0.0007576	0.0935479	-0.1214028	-0.1896307	0.072464
## do	-0.0302646	-0.0727378	-0.0926549	-0.0626465	-0.1555949	0.036756
## is	0.1336907	0.0797518	-0.1785689	0.0632460	-0.1423981	-0.053693
## or	-0.1089584	-0.1039416	-0.0132435	0.0042732	0.0412992	-0.177685
## this	0.2334027	0.0630079	-0.0218351	-0.0906833	-0.1114024	0.034180
## all	0.0844019	0.0628760	0.0426171	0.1138584	0.1803859	-0.052679
## down	0.0292295	-0.0751349	-0.0179187	0.0165084	0.0005723	-0.140894
## it	-0.0292278	-0.0288241	-0.2022608	-0.2491168	0.0007965	0.065907
## our	-0.1253024	-0.0447230	0.0946293	-0.0059994	0.1822235	-0.004053
## to	0.1133667	-0.1822340	-0.0075929	-0.1504039	0.1180770	0.042107
## also	-0.1895352	0.0803221	0.0144453	0.1038795	-0.2159480	-0.021350
## even	0.0412006	-0.0053735	0.1988205	-0.1080156	-0.0070907	-0.085335
## its	0.0925980	-0.1364080	-0.1003783	0.0562824	0.1573799	-0.087837
## shall	0.0275950	0.0069202	-0.0851937	0.0941129	-0.2189935	0.061184
## up	0.0440077	-0.0079942	0.1198425	-0.1167414	0.0746236	-0.092781
## an	0.1987379	0.0008641	0.0911709	-0.1070558	-0.0145571	0.066112
## every	0.1006472	0.0201706	-0.1635864	0.1268195	-0.1723828	-0.016974
## may	-0.0002659	-0.1579650	-0.1937055	0.0956668	-0.1181080	0.055328
## should	-0.0633278	0.0130290	-0.0238200	-0.2547560	0.0931800	-0.055491
## upon	0.2414271	-0.1504749	0.0467621	-0.1254395	0.0574721	0.071282
## and	-0.2969149	0.0829995	-0.0079523	0.0648104	0.1141941	-0.069459


Example

What variables matter most?

##	rowname	PC1	PC2	PC3	PC4	PC5	PC6
## 1	and	-0.2969149	0.0829995	-0.0079523	0.0648104	0.1141941	-0.069459
## 2	of	0.2532369	0.0465962	0.1820845	0.0849970	-0.0677631	-0.024695
## 3	upon	0.2414271	-0.1504749	0.0467621	-0.1254395	0.0574721	0.071282
## 4	this	0.2334027	0.0630079	-0.0218351	-0.0906833	-0.1114024	0.034180
## 5	one	-0.2310547	-0.0273376	0.0517796	0.0560865	-0.2252644	-0.001185
## 6	more	-0.2192323	-0.1299366	0.1103656	0.1446700	0.0100714	0.140034
## 7	their	-0.2098190	0.0526373	-0.0558632	-0.1481704	0.1857935	0.126837
## 8	an	0.1987379	0.0008641	0.0911709	-0.1070558	-0.0145571	0.066112
## 9	a	0.1981759	-0.0007576	0.0935479	-0.1214028	-0.1896307	0.072464
## 10	the	0.1959157	0.0652826	-0.0653657	0.0971773	-0.0635347	-0.142684
## 11	also	-0.1895352	0.0803221	0.0144453	0.1038795	-0.2159480	-0.021350
## 12	into	-0.1646578	-0.0768915	0.0947230	0.1017412	-0.0845516	0.010113
## 13	there	0.1570871	-0.2135796	0.0912550	-0.0034868	-0.0925714	0.003078
## 14	in.	0.1517917	0.0252016	-0.0993965	-0.1591765	-0.1155705	0.118976
## 15	which	0.1406579	0.0833763	-0.0633439	0.1863356	0.1818572	-0.010853
## 16	been	0.1397300	0.2585870	-0.0832399	0.0552028	-0.0521795	0.108904
## 17	is	0.1336907	0.0797518	-0.1785689	0.0632460	-0.1423981	-0.053693
## 18	than	-0.1292803	-0.2098651	0.0407369	0.1246189	-0.1711510	0.261084
## 19	our	-0.1253024	-0.0447230	0.0946293	-0.0059994	0.1822235	-0.004053
## 20	only	-0.1252312	0.0404905	-0.1818237	-0.1444484	-0.0538031	-0.062422
## 21	with	-0.1202768	-0.0353634	0.1731743	0.1397497	-0.0882275	-0.029550

Example

This doesn't really help us visualize the data...




Example

Locations of observations on new axes:

##	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8
## 1	0.255883	-2.75723	-0.84061	4.22604	3.580825	7.87903	-0.74834	-2.742008
## 2	-5.333328	4.41925	0.10285	-1.85960	0.385100	2.77020	-2.05465	0.310339
## 3	-7.933306	-1.90187	-2.06613	0.17152	0.444284	-0.88954	2.73854	1.482792
## 4	-10.132733	-3.01485	-0.76659	0.22149	0.388497	-0.82803	-1.44149	2.264506
## 5	-8.922169	-2.22472	3.78289	-0.08058	-0.187915	0.18637	-3.19444	-4.563181
## 6	-0.664084	3.30507	3.19383	1.00454	1.497993	1.68512	0.78693	2.868702
## 7	1.140376	1.88974	5.25083	-1.66296	1.699322	-0.41244	-0.62420	-0.832843
## 8	0.358401	1.23172	2.36159	0.81826	1.955048	-1.53049	-1.06244	1.153950
## 9	0.140885	4.29456	2.16608	0.56263	-0.004672	0.68167	0.30781	1.427079
## 10	-2.035723	-0.31634	-1.02594	1.82065	1.242992	-0.64046	0.86073	-0.908341
## 11	-0.637666	-1.45658	3.78830	0.35541	-0.251304	-0.36100	-3.51373	0.569111
## 12	1.392538	-0.17138	1.80856	1.19238	1.713521	-1.00059	-2.12927	1.822116
## 13	-2.132165	-6.97474	4.04474	4.79525	-6.312285	0.88882	-2.32670	1.957892
## 14	-1.017369	1.67820	-1.35167	2.20414	2.951174	1.91323	-0.15203	-1.600599
## 15	2.746062	0.09324	1.33704	2.02277	2.416346	0.17314	-2.49471	2.170778
## 16	1.346831	-1.09305	2.17638	-2.53938	0.597100	-0.88129	-0.16961	-0.264660
## 17	0.945306	0.21010	4.27734	2.09572	2.968605	0.35147	1.51480	-1.561375
## 18	3.014215	-0.18193	0.45892	1.34565	-1.801528	-1.42809	-1.18131	2.533578
## 19	1.785756	0.37943	0.92432	-0.41027	0.024650	-0.17518	-1.19383	0.674022
## 20	1.699876	-1.45300	-1.79566	1.69022	3.632504	-1.03967	0.60535	1.341652
## 21	0.688843	0.02942	2.43645	-3.92072	0.986641	-0.41756	-2.71591	-0.324139


```
new_dims_df %>%  
  ggplot(aes(x = PC1, y = PC2, color = auths_known)) +  
  geom_point()
```



Example

Standard deviations of PC scores:

```
## [1] 2.584293976481497168 2.224602605692588053 2.133495077907519910
## [4] 2.009679212438630014 1.771199010414047370 1.737846153429807972
## [7] 1.658192994671916720 1.593432897204484666 1.567765465416049331
## [10] 1.498153968179446682 1.448627829241023512 1.383202255873344555
## [13] 1.348835290469535098 1.329864883494040306 1.290075723747528214
## [16] 1.263656822413196323 1.211546834748383983 1.179468039028208848
## [19] 1.140273064226865252 1.136685249904107353 1.091881595973659413
## [22] 1.024359742653421224 1.002539216556411761 0.988309265374064272
## [25] 0.966758063464953188 0.953955378557446387 0.913371285706331415
## [28] 0.879219144916352668 0.866397791514436300 0.835362796306243660
## [31] 0.813744862497518318 0.795836559214300077 0.777288905706550737
## [34] 0.755507991799209577 0.729325759585313982 0.723553021174748179
## [37] 0.668529227774796841 0.644640218695886280 0.608063670676834200
## [40] 0.599862980594791706 0.589869259793525114 0.549984140873157168
## [43] 0.518576401942239196 0.490365964788197883 0.485951132126856200
## [46] 0.473420966039042657 0.447648096786177840 0.410247671033032724
## [49] 0.398219462937036084 0.381729038085061145 0.356072979106936693
## [52] 0.322926868470927830 0.302289519970545262 0.278418759166978558
## [55] 0.254888335158211798 0.246279930847948220 0.227060020169778221
## [58] 0.209288942626959323 0.201384112256443981 0.181411412891078455
## [61] 0.162645660545995191 0.152090980040969159 0.125916504292499787
## [64] 0.109322167263386283 0.082184427523104023 0.065913614423482467
```

Example

Cumulative variances:

```
cumul_vars <- cumsum(pc$sdev^2)/sum(pc$sdev^2)
cumul_vars
```

```
## [1] 0.09541 0.16611 0.23113 0.28883 0.33365 0.37679 0.41607 0.45234 0.48745
## [10] 0.51952 0.54950 0.57683 0.60282 0.62809 0.65186 0.67467 0.69564 0.71552
## [19] 0.73409 0.75255 0.76958 0.78457 0.79893 0.81288 0.82623 0.83923 0.85115
## [28] 0.86219 0.87292 0.88289 0.89235 0.90139 0.91003 0.91818 0.92578 0.93326
## [37] 0.93964 0.94558 0.95086 0.95600 0.96097 0.96529 0.96914 0.97257 0.97594
## [46] 0.97915 0.98201 0.98441 0.98668 0.98876 0.99057 0.99206 0.99337 0.99447
## [55] 0.99540 0.99627 0.99700 0.99763 0.99821 0.99868 0.99906 0.99939 0.99961
## [64] 0.99979 0.99988 0.99994 0.99997 1.00000 1.00000 1.00000
```

Example

```
plot(cumul_vars)
```

Details

How many PCs should we use?

No single answer; people often do "enough for 90% variance covered" or similar.

How do you do modeling with PCA?

You don't. It's a data *preprocessing* step.

You would then use PC1, PC2, etc **instead of** your original predictors.

OR you might use the variable importance measures ("pc loadings") to help decide which predictors to keep.

Pros and Cons

Pros:

- Reduces dimension while still letting all original predictors be "involved"
- Computationally fast for big data
- Axis rotations are interpretable!
- Dropping off lower PCs gets rid of noise (maybe)

Cons:

- Using PCs in interpretable models makes them uninterpretable. (What does the coefficient of PC1 mean in real life?)
- No magic answer for how many PCs to use.
- Dropping off lower PCs gets rid of useful info (maybe)

But wait - isn't this LDA?

Recall that **LDA** also found *scores* for your observations based on a *linear combination* of the original predictors.

So what's the difference?

LDA loadings are trying to maximize the **difference in mean scores across categories**.

(*supervised* method)

PCA loadings are trying to maximize the **variance of the scores** on the PC axes.

(*unsupervised* method)

Try it!

Open **Activity-PCA.Rmd**

Apply PCA to the cannabis data

Interpret the PC rotations

Plot the data on the first two axes, and color by Type.

Choose a "good" number of PCs to use.

Fit a KNN classifier using only your chosen PCs. How does the accuracy compare to when you use all the original predictors?